

BRRTRouter: Coroutine-Powered OpenAPI Router for Rust

BRRTRouter is a high-performance request routing library for Rust, powered by stackful coroutines (via the **May** crate) instead of `async` / `await`. It allows you to define a REST API **entirely from an OpenAPI 3.1 specification**, and wires up request handlers based on that spec ¹. BRRTRouter aims to combine the ergonomics of dynamic API frameworks with the performance of Rust, without relying on Tokio or any asynchronous runtime.

Overview and Features

- **OpenAPI-Driven Routing:** You supply an OpenAPI 3.1 document (YAML/JSON) describing your endpoints, schemas, and operations. BRRTRouter will parse this spec (using `load_spec`) and produce an internal route table (paths, methods, parameters, responses, etc.) ². This means your API routes are **defined declaratively** by the spec, avoiding duplicate manual routing code.
- **Coroutine-Based Concurrency:** Under the hood, BRRTRouter uses the **May** coroutine library instead of Rust's `async` system. May provides lightweight user-space threads (M:N scheduling) that run on a pool of OS threads. Each incoming HTTP request is handled in a separate coroutine, allowing high concurrency without OS thread explosion. The included HTTP server is built on `may_minihhttp`, a port of Tokio's minihhttp example to May ³ ⁴. This lets BRRTRouter handle many connections with minimal overhead, while you write handler code in a blocking style (no `.await` necessary).
- **Automatic Request Dispatch:** BRRTRouter maps incoming requests to the correct handler function based on the path and HTTP verb using the parsed OpenAPI spec. It constructs a `Router` (which matches paths to operation identifiers) and a `Dispatcher` (which holds the actual handler functions/objects) ⁵ ⁶. When a request comes in, the Router finds the route and the Dispatcher invokes the corresponding handler.
- **Typed Handlers and Auto-Deserialization:** You can register handlers with strong typing for request data. For example, if an operation expects a JSON body of a certain schema, you can define a Rust struct for it and register a handler using `dispatcher.register_typed("operationId", MyHandlerStruct)`. BRRTRouter will automatically deserialize the JSON body into your struct and pass it to the handler. This provides compile-time type checking for request data. In the included Pet Store example, a code-generated registry uses `register_from_spec` to register each operation's handler with the Dispatcher (parsing bodies into typed structs) ⁷. Handlers produce a `HandlerResponse` (status code, headers, body) which BRRTRouter will serialize to JSON as needed.
- **Middleware and Extensibility:** The library includes a middleware system (extended in recent commits) to enable cross-cutting concerns like logging, tracing, or metrics ⁸. For example, a `TracingMiddleware` was introduced to track coroutine stack usage and OpenTelemetry spans in

tests ⁸. Middlewares can wrap around handlers to intercept requests/responses. A metrics endpoint and a health-check endpoint have been added as built-in features ⁹ ¹⁰, and static file serving is supported via a “StaticFiles” loader ¹¹.

- **Server-Sent Events (SSE) Support:** BRRTRouter supports SSE for streaming data to clients. It includes an SSE channel implementation and tests for server-sent events ¹². (For example, a handler can push events to a client over a long-lived HTTP connection.) The SSE implementation required some special handling of HTTP headers (e.g., proper `Transfer-Encoding`) which has been fixed in the parser ¹³. The library ensures non-blocking writes for SSE so that one slow client doesn't stall others.
- **Hot-Reloadable Routes:** You can enable **hot-reloading** of the OpenAPI spec at runtime, ideal for development. A `watch_spec` function is provided which watches the spec file for changes and reloads the routes dynamically ¹⁴ ¹⁵. When the spec file is modified on disk, `watch_spec` will re-parse it and update the in-memory Router (guarded by an `Arc<RwLock<Router>>`). It also invokes a user-provided callback so you can update your Dispatcher or other state accordingly ¹⁶ ¹⁷. This means you can add new endpoints to the spec and see them live without restarting the server. The router updating is done under a write-lock to ensure thread safety while existing requests continue using the old routes until the swap is complete ¹⁵.
- **High Performance:** By using May's coroutine scheduler and avoiding the overhead of futures and `await`, BRRTRouter strives for excellent performance comparable to or exceeding traditional async frameworks. The coroutine model can make context switches cheaper and allow a more imperative programming style. The maintainers have benchmarked May's HTTP handling in the TechEmpower Web Frameworks benchmark in the past (via `may_minihhttp`), which showed competitive throughput. The stack size for coroutines is configurable; BRRTRouter sets a custom stack size (e.g. 32KB in tests) to balance memory usage and avoid overflow ¹⁸. *(In fact, the test suite uses a trick of setting an odd stack size to make the runtime print stack usage stats on drop for debugging ¹⁹.)*

Usage Example

Below is a simplified example of how you might use BRRTRouter in a Rust application:

1. **Prepare OpenAPI Spec:** Write your OpenAPI 3.1 spec (YAML or JSON) describing paths, operations, parameters, and schemas. For instance, `examples/openapi.yaml` in this repo defines a simple Pet Store API.
2. **Load Routes from Spec:** Use `brtrrouter::load_spec(path)` to parse the spec into a vector of routes (internally `RouteMeta` objects). Then create a `Router` from these routes to enable path matching:

```
let (routes, _errors) = brtrrouter::load_spec("examples/openapi.yaml").unwrap();
let router = Router::new(routes.clone());
```

This will produce route matchers for all paths and HTTP methods defined in the spec ². You can share the `Router` across threads/coroutines by wrapping it in an `Arc<RwLock<Router>>` if you plan to hot-reload or modify it at runtime ²⁰.

1. **Register Handlers:** Create a `Dispatcher` and register a handler for each operation (usually identified by the operationId or an internally generated key). There are two ways:
2. *Raw Handlers:* Use `dispatcher.register_handler("operationName", |req: HandlerRequest| { ... })` to register an inline closure or function. You'll manually handle parsing the request data from `req` and sending a response.
3. *Typed Handlers:* Define a struct or type that implements the `Handler` trait (with an associated request type for automatic deserialization). Then use `dispatcher.register_typed("operationName", MyHandlerType)`. `BRRTRouter` will deserialize the JSON body into `MyHandlerType::Request` and call your handler. This method is safer and less error-prone.

For example, in the Pet Store demo, a code-generated module `pet_store::registry` provides a function to register all spec-defined handlers at once:

```
let mut dispatcher = Dispatcher::new();
unsafe { pet_store::registry::register_from_spec(&mut dispatcher, &routes); }
```

This `register_from_spec` call will iterate over the routes and link each one with a concrete handler (the handlers could be auto-generated stubs or real implementations you filled in) ⁶ ²¹. The `unsafe` here is used because the codegen might be doing some tricks to cast generic handlers – this is an area planned to be improved with safer abstractions.

1. **Start the Server:** `BRRTRouter` provides an integration with `may_minhttp` to run an HTTP server. You create an `AppService` with your `Router` and `Dispatcher`, then start the server:

```
use brrouter::server::AppService;
use may_minhttp::HttpServer;
use std::sync::{Arc, RwLock};
use std::net::SocketAddr;

let service = AppService::new(
    Arc::new(RwLock::new(router)),           // thread-safe router
    Arc::new(RwLock::new(dispatcher)),       // thread-safe dispatcher

    std::collections::HashMap::new()         // (optional) shared state, e.g. SSE
);
let addr: SocketAddr = "127.0.0.1:8000".parse().unwrap();
HttpServer(service).start(addr).unwrap();
```

This will spawn the HTTP listener on the given address and start accepting connections. Each incoming request is handled by invoking your `service`, which in turn uses the Router/Dispatcher to produce a response. The `HashMap::new()` passed to `AppService::new` is a placeholder for any application state you want to share (for example, it could hold active SSE client channels, or other global data accessible to handlers).

Note: In tests, the server is often started on an ephemeral port (by binding to `127.0.0.1:0`) and a small delay is used to ensure the background threads are listening ²² ⁴. In a real application, `HttpServer.start()` is synchronous and will return a handle or error immediately. You can call `.join()` on the returned handle to block on the server (or drop it to run server in the background).

1. **Make Requests:** With the server running, you can use any HTTP client or tool (cURL, Postman, etc.) to call your API. `BRRTRouter` will route the request to the correct handler. For example, a `GET /pets/42` request will be dispatched to the handler for the “GetPet” operation, and a `POST /pets` will go to the “CreatePet” handler, as defined in the OpenAPI spec.
2. **Hot Reload (Optional):** During development, you can enable hot-reload so that editing the spec file updates the routes live. For example:

```
use brtrouter::spec::watch_spec;
use notify::RecommendedWatcher;
let router_arc = Arc::new(RwLock::new(router));
let watcher: RecommendedWatcher = watch_spec("examples/openapi.yaml",
router_arc.clone(), |new_routes| {
    // Re-register handlers or update dispatcher if needed
    dispatcher.update_routes(&new_routes);
}).expect("failed to watch spec");
```

Now if you edit the spec file, the router will be refreshed in-memory. The callback gives you the new routes so you can adjust the dispatcher – for instance, you might auto-register any new endpoints or log the changes. The watch uses the `notify` crate internally and filters for file modify/create events ¹⁴. When triggered, it reloads the spec and replaces the router inside the `Arc<RwLock<Router>>` with a new one ²³ ¹⁵. Existing requests continue with the old router until they finish, and new requests will use the updated routes. This feature is especially useful for iterative development, though in production you’d likely keep the routes static (or restart to deploy API spec changes).

Internals and Architecture

`BRRTRouter`’s core components include:

- **Spec Parser (`spec` module):** Handles reading the OpenAPI YAML/JSON and converting it into internal route representations. It ensures path templates, parameters, response schemas, etc., are understood. If the spec has unsupported features or errors, those are returned alongside the routes. The parser also marks which routes should be handled as JSON, which as streaming (SSE), etc., based on spec extensions or content types.

- **Router (router module):** A structure that holds route matchers for each endpoint. It likely uses a trie or hash map internally to map an incoming `method + path` to a route identifier or metadata (the code is split across submodules after a refactor ²⁴). The Router can perform dynamic path parameter extraction. For example, if a path `/pets/{id}` is defined, the Router will match `/pets/42` and extract the `id`. The result of a lookup is a `RouteMeta` containing the operation name (or an identifier) and any extracted params.
- **Dispatcher (dispatcher module):** This maps operation names or route IDs to actual handler logic. A handler can be a closure, a function, or an object implementing a `Handler` trait. The Dispatcher also manages **typed deserialization**: if a route expects a certain JSON request body, the Dispatcher will parse the `HandlerRequest` body into the appropriate Rust type before invoking the handler. It uses `serde` under the hood (the request structs derive `Deserialize` ²⁵). The Dispatcher is also responsible for sending back the response via the request's reply channel. In the internals, each `HandlerRequest` carries a oneshot channel (`reply_tx`) to send a `HandlerResponse` back to the HTTP server loop ²⁶ ²⁷. This decouples the request-handling coroutine from the network sending logic.
- **AppService (server module):** This ties the Router and Dispatcher together into a callable service that the HTTP server can use. It implements the interface expected by `may_minihhttp`: essentially a function `service(request) -> response`. When an HTTP request comes in:
 - The service looks up the route via Router (under a read-lock if hot-reload is enabled) and obtains the route metadata.
 - It packages the HTTP request data (path, query params, headers, body, etc.) into a `HandlerRequest` object.
 - It hands off the `HandlerRequest` to the Dispatcher, which finds the corresponding handler and executes it (possibly in a new coroutine if the design dictates).
 - The handler produces a `HandlerResponse` (which includes an HTTP status code and possibly a JSON body or stream).
 - The response is sent back to the client via the `may_minihhttp` connection.

The AppService may also apply global middleware. For example, a tracing middleware can wrap the request handling to record how much stack was used or log the request path ⁸. Because everything runs on May coroutines, you can use blocking operations in your handlers (such as database calls or file I/O) – just be aware that a truly blocking call (that doesn't yield to the scheduler) could block other coroutines on the same OS thread. Many standard operations (network, sleep, etc.) are patched to be non-blocking in May's runtime, but long CPU-bound tasks won't yield by default. You should manually yield or spawn a new coroutine for such tasks to keep the system responsive.

- **Shared State and SSE:** The third parameter to `AppService::new` is a `HashMap` that can hold any shared state needed by handlers or middleware ²⁸. This could be used to store things like database connection pools, caches, or SSE broadcast channels. In particular, SSE handlers might insert each client's channel into this map so that other parts of the app can push events to them. The design allows passing this shared state down to handlers (likely via the `HandlerRequest` or a global). The current implementation uses a simple `HashMap`, but this is an area that might be generalized to a proper state container trait in the future.

- **Error Handling:** If a handler returns an error or panics, ideally BRRTRouter should catch that and convert it into a 500 Internal Server Error response, rather than crashing the whole server. **However, as of now there are some gaps in error handling for panics.** The test suite identified that if a handler panics, the framework did not properly catch it, causing the test itself to panic or hang ²⁹ ³⁰. A TODO exists to fix this by wrapping handler calls in `std::panic::catch_unwind` and then responding with a 500 code ³¹ ³². This is a known area for improvement (see “*Improving Reliability*” below).
- **Security and Authorization:** OpenAPI allows defining multiple security schemes for an operation (e.g., an endpoint might be accessible via either an API key or a JWT token). BRRTRouter supports multiple security providers and will attempt each one to authenticate a request. The implementation for this spawns separate coroutines for each security provider to check credentials concurrently, and if one succeeds, the others are canceled. This design can reduce latency (whichever auth method returns first wins). **However, there is a known timing issue here:** The test for multiple security providers occasionally fails due to a race condition in canceling the losing coroutine ³³. Essentially, one auth provider may finish just as the other is also yielding a result, and if not synchronized, it can lead to two responses or other undefined behavior. This is another aspect to be refined (either by making the security checks sequential or introducing a more robust synchronization so that exactly one provider’s result is used).

Test Suite Status and Flakiness Analysis

BRRTRouter comes with an extensive test suite covering routing, dispatching, SSE, hot reload, etc. Because of the concurrent nature of the library, some tests have been **flaky (intermittent failures)** in CI. The developers have marked certain tests with `#[ignore]` or added TODO comments to diagnose issues:

- **Panic Recovery Tests:** There are tests intended to verify that if a handler panics, the server responds with HTTP 500 and continues running. Currently, these tests are ignored because the panic is not being caught correctly, causing the test itself to panic (and thus fail) ²⁹. For example, `test_panic_handler_returns_500` and `test_panic_recovery` are annotated with `#[ignore]` until the panic-catching logic is implemented ³¹ ³². The root cause is simply that the `AppService`/Dispatcher does not yet wrap handler calls in a `catch_unwind`. **Planned fix:** Add panic handling around handler invocation, and verify these tests pass (i.e., the response status should be 500 and the process should continue). This will allow re-enabling those tests.
- **Multiple Security Providers Race:** As mentioned, `test_multiple_security_providers` sometimes fails. A comment in the test notes it “fails intermittently due to timing issues with the coroutine cancellation” ³³. The pattern observed in CI logs was that occasionally the wrong status code or response would be seen, likely because one auth coroutine was not properly canceled in time. **Planned fix:** Revise the security provider handling to ensure deterministic outcome. This could involve using a synchronization primitive (e.g., an `AtomicBool` flag to indicate “auth succeeded” that other coroutines check before sending a response) or simpler, executing providers in sequence (since usually there are at most two or three, the performance impact would be minimal). Until this is fixed, this test might remain flaky.

- **Stack Size and Overflow:** Early on, some tests encountered stack overflows or aborts when running many coroutines due to stack exhaustion. By default, May assigns a fixed stack size to coroutines (2 MB unless changed). The project maintainers discovered that certain deeply nested operations (or perhaps large JSON parsing) could use a lot of stack. To catch any issues, they configured all test coroutines to use a smaller stack (32 KB) and even leveraged May's debug feature: setting an odd stack size (e.g., 0x8001) causes the runtime to print out stack usage statistics on coroutine exit ³⁴ ¹⁹. They added `may::config().set_stack_size(0x8000)` at the start of every test to standardize the stack allocation ³⁵ ³⁶. This ensures tests don't randomly fail on machines with different default stack settings, and it flushes out any function that might blow the 32KB budget. So far, with 32KB stacks, the tests pass reliably (no more overflows after they fixed one bug in SSE header parsing that was causing unbounded recursion).
- **Hot Reload Test:** There is an integration test for the hot-reload functionality (`test_watch_spec`). This test spawns a background thread to watch the spec file, then makes a change to the file and waits to see if the router updates. Timing issues could occur (file notifications might take a moment). To make the test robust, it waits for a specific trace event indicating the reload happened ³⁷ ³⁸. The project introduced a tracing mechanism (`tracing.wait_for_span("secret")`) in the security tests, and presumably similar in reload tests) to coordinate the test flow ³⁹. The test logs on CI showed that waiting for these spans greatly reduced flakiness by ensuring the test only proceeds after the background work is done. This pattern of using instrumentation in tests is a notable approach in BRRTRouter.
- **Server Start/Stop Races:** In some tests, the server is started in a background coroutine and the test then immediately begins sending requests. A small sleep (50 ms) was inserted after starting the server ⁴ to give the listener time to actually bind and accept. If the CI machine is slow, there's a theoretical possibility that 50 ms isn't enough, causing a connection failure. While this hasn't been a major source of flakiness, a more bulletproof approach would be to have the server signal readiness (for example, by sending on a channel once it's listening) instead of using a fixed sleep. Implementing a handshake in `AppService` (or extending `HttpServer` to return once ready) would make tests start reliably. This is a minor improvement to consider for test robustness.

In summary, the intermittent test failures primarily stem from **race conditions in concurrent code and missing error handling**. The team has been actively identifying these via CI logs and adding instrumentation to diagnose them. For instance, they temporarily added extensive tracing (with a fake OpenTelemetry collector) to the coroutine tests in order to track down a sporadic failure ⁸. That change was later reverted once enough data was gathered ⁴⁰ ⁴¹. Going forward, each flaky test has a clear action item (catch panics, fix auth race, improve sync), so we expect the suite to stabilize.

Improving Reliability and Concurrency Handling

Based on the analysis above, here are recommendations to harden BRRTRouter's codebase:

- **Implement Panic Recovery:** Wrap handler execution in `std::panic::catch_unwind` inside the Dispatcher or AppService. This way, if a handler panics, the coroutine can catch it, log it, and return an HTTP 500 instead of propagating the panic. This will prevent a rogue handler from tearing down the whole server and will allow re-enabling the ignored panic tests ²⁹. Care must be taken to drop

or reset any state (e.g., lock guards) if a panic occurred. Since May's coroutines will unwind just like threads, catching unwind at the top of each request coroutine is the safe approach.

- **Fix Multi-Auth Race:** Refactor the multiple security provider logic to ensure only one authentication result is ever used. One approach is to run them sequentially (try first method, if it fails, then try next) which is simpler and nearly as fast for expected scenarios. If parallel execution is desired for performance, use an atomic flag or synchronization: for example, spawn both checks, but have each check race to an `Arc<AtomicBool>` flag; whichever sets it to true first "wins" and proceeds to use its result, while the other checks the flag and exits if it's already true. Additionally, using a bounded channel for replies could ensure that only the first reply is received. The goal is to remove timing sensitivity so the test passes consistently. After this fix, remove or update the TODO in `test_multiple_security_providers` ³³.
- **Deterministic Server Startup/Shutdown:** Implement a mechanism for graceful startup and teardown of the HTTP server in tests. For example, modify `HttpServer::start()` to return a handle that has a method `ready() -> bool` or send a signal on a channel once the listening socket is active. For shutdown, provide a way to stop the server loop (maybe by closing the listening socket or using an atomic quit flag that causes the server coroutine to break out). Currently, tests simply drop the join handle, which in May's runtime likely detaches the coroutine to run until program end. This is generally fine (the OS will clean up the thread when the test process exits), but it's not elegant. A controlled shutdown would allow cleaning up between tests if needed and would be essential for a long-running application that needs to reload or stop the server on command.
- **Increase Test Determinism:** Continue using tracing spans or other signals to coordinate test threads with background tasks. This has proven effective. For instance, after triggering a file change in the hot-reload test, wait for a "spec reloaded" log or span before asserting the new routes are in place. This eliminates race conditions in tests where timing can vary. Such instrumentation can be behind a test-only feature flag so it doesn't affect release performance.
- **Validate Stack Usage in Critical Paths:** While the reduced coroutine stack size (32KB) is used in tests, in production one might use a larger stack to be safe. It would be good to document the expected stack needs or even perform runtime checks (if May had an API for current usage, which it does not currently expose ³⁴). One idea is to have an integration test that hits a deep recursion (or large JSON) scenario to ensure it doesn't overflow the chosen stack. Given that May's default is around 2MB, the 32KB in tests is quite small, so if all tests pass under that, production use with default stack should have a big safety margin. Still, if any handler requires significantly more stack (e.g., very deep JSON structures with recursive parsing), it should be converted to an iterative approach or documented.
- **Race-Free Hot Reload:** The current hot-reload implementation using `notify` is robust, but tests should also ensure that no requests are being processed using a partially updated router (the use of `RwLock` prevents this by making the swap atomic ¹⁵). It might be worth adding a small delay or check after a reload to ensure the new routes are live. Also, on reload, the `on_reload` callback should carefully update any global state. For example, if routes were removed, the Dispatcher should remove those handlers to avoid memory leaks. These edge cases could be tested more thoroughly.

- **Concurrency Stress Testing:** To further harden the code, one could write a stress test that fires a large number of concurrent requests (perhaps using multiple threads or processes) at the server to ensure that the scheduling in May holds up and there are no data races. The current tests cover functionality but not high concurrency scenarios. Using tools like loom (for checking concurrency invariants) is challenging with foreign crates like May, but a manual stress test can still be valuable. For instance, concurrently exercise hot-reload while serving requests, or issue parallel requests that intentionally cause handler panics, to see if the system remains stable.
- **Upgrade May and Dependencies:** Ensure the `may` and `may_minhttp` dependencies are up to date with the latest fixes. The May crate (by Xudong Huang) is mature but not as widely used as Tokio, so keeping an eye on its issue tracker for any bug reports (especially around unwinding, I/O, or synchronization) is prudent. If any issues are found (e.g., a bug in channel implementation or a corner-case in coroutine cancellation), consider contributing a fix or workaround. So far, BRRTRouter's usage of May has uncovered the stack usage quirk (addressed via config) and the need for better panic catching (which is in BRRTRouter's purview to implement).
- **Documentation of Guarantees:** From a reliability standpoint, clearly document what BRRTRouter guarantees (and what it doesn't yet). For example: Does it guarantee order of middleware execution? (Yes, presumably FIFO order as added). Does it catch panics in handlers? (Not yet, but will soon). Are handlers allowed to spawn their own coroutines (yes, they can spawn via May's API if they need to perform parallel subtasks). What happens if a handler fails to send a response on the reply channel? (Likely the connection will hang; maybe enforce a timeout or at least log an error). Defining these behaviors will guide future improvements and user expectations.

Performance Considerations

BRRTRouter is designed for high performance, but here are some notes and tips:

- **Coroutine Overhead:** May's coroutines are very lightweight (on the order of microseconds to spawn) and use a work-stealing scheduler across threads. This allows tens of thousands of concurrent coroutines. The elimination of `async/await` state machines can reduce per-request overhead. However, context switching is not free – if your handlers are purely CPU-bound, you could still saturate a core. In such cases, adding more OS threads (via May's configuration) can help utilize more cores. Profile the application to find bottlenecks. The good news is that for I/O-heavy workloads (common in web servers), coroutines can park and yield control efficiently.
- **I/O and Blocking Calls:** Under the hood, `may_minhttp` uses non-blocking sockets and an event loop, so network I/O will not block an OS thread. If your handler performs other blocking syscalls (like reading a file, interacting with a database over a blocking client), those will block the entire OS thread's coroutine scheduler. To avoid this, prefer asynchronous versions (if available) or spawn a separate thread for truly blocking operations. You can also increase the thread pool size for May so that other requests aren't stalled by one blocking call. Future versions of BRRTRouter might integrate with asynchronous backends or provide hooks to offload blocking tasks.
- **Memory Footprint:** Each coroutine has its own stack. The default stack size (2MB) multiplied by thousands of coroutines could be heavy on memory. In practice, most coroutines won't use the full

2MB. You can tune the stack size based on your endpoint requirements. For example, if you know none of your handlers use deep recursion or huge stack allocations, you might set a smaller stack size globally (as done in tests with 32KB) to save memory. On the other hand, if you do JSON deserialization of very large objects, ensure the stack is large enough or switch those to heap allocations. The project's approach of using a moderate stack and monitoring usage is a good guideline.

- **Routing Speed:** Route lookup in the Router is designed to be fast (likely $O(\text{depth of path})$ or $O(1)$ with hashing). The library could be optimized further by using techniques like HTTP method partitioning or precomputing small finite state machines for path patterns. As of now, performance hasn't shown an issue, but if you have a very large number of routes (hundreds or thousands), keep an eye on routing latency. A possible future enhancement is to use a prefix tree for routes to speed up matching of long common prefixes.
- **Avoiding Lock Contention:** With hot-reload enabled, the Router is behind an `RwLock`. Reads vastly outnumber writes in typical usage, so this should not be a bottleneck (reads are simultaneous as long as no write is happening). Writes (spec reloads) are infrequent. If you disable hot-reload in production, you can potentially avoid the lock entirely by using the Router directly (since it's immutable after construction). The Dispatcher is not locked per request in the current design – handlers are looked up by an internal map, and each request's `HandlerRequest/Response` channel is separate, so there's no global mutex on sending responses aside from any synchronization in the channel itself (which is per-request). This means the architecture is *mostly lock-free* in the steady state, which is good for throughput.
- **Metrics and Monitoring:** The library added a **metrics endpoint** in a recent update ⁹. This likely provides basic stats (perhaps number of requests served, uptime, etc.). Utilize it to monitor performance. In addition, consider using the tracing integration – you could attach an OpenTelemetry exporter to gather metrics like response times, coroutine schedules, etc. In tests, a custom tracing collector was used to check that certain spans (like `"secret"` for auth events) occurred ³⁷. Similar hooks can be used in production to gather telemetry.

Future Roadmap

BRRTRouter is still evolving. Based on the repository issue tracker and author comments, here are some expected improvements and planned features:

- **Comprehensive Validation:** Implementing request **validation** against the OpenAPI schema is a key goal (mentioned in issue #5, *"Validation & Route handler & controller Roadmap"*). This means if your spec says a query parameter must be an integer between 1 and 100, the framework should automatically check that and reject out-of-range values with a 400 Bad Request before your handler is called. Currently, basic parsing is done, but rich validation (formats, string lengths, etc.) may be limited. Expect more automatic validation of inputs and outputs according to the OpenAPI spec, reducing the amount of manual checking needed in handlers.
- **Code Generation for Types and Handlers:** The project has a code generation component (the `templates/` directory and the `photon` framework are hints to this ⁴²). The goal is to generate

Rust types for schemas defined in the OpenAPI spec and skeleton handler functions that developers can fill in. In a recent forum post, the author mentioned they had “implemented code generation for routes and handlers” and the next step was generating typed handlers and structs, followed by hot reload integration ⁴³. This likely means a future version will allow you to run a codegen tool that reads the OpenAPI spec and produces a Rust module (like the `pet_store` module in the example) with:

- Data transfer object structs for request and response bodies (with `serde::Deserialize/Serialize` derived).
- An enum or constants for each operationId.
- Stub functions or a trait for each operation that you can implement, so that linking spec to code is mostly automatic.

This will simplify usage: instead of manually writing handler registration code, you'd just implement trait methods or fill in the generated functions for each endpoint.

- **Better Developer Experience (DX):** The wrapper framework **Photon** (mentioned by the author ⁴²) is likely to provide a higher-level API on top of BRRTRouter. Photon might incorporate features like dependency injection of the shared state, macros to start the server, and perhaps integration with database or auth libraries, making it feel more like familiar web frameworks (but built on the coroutine paradigm). As of the latest update, Photon is just a placeholder, but it indicates the direction for ergonomics improvements.

- **Stability and Releases:** The author has indicated an intention to **publish the crate** to crates.io after ironing out the remaining issues ⁴³. Before a 1.0 release, expect the following to be addressed:

- All tests reliably passing (no ignores or flakes).
- Complete documentation (a thorough README – possibly this document – and a RustDoc or book with guides).
- Audit of unsafe code (e.g., the `unsafe { register_from_spec(...) }` should be well-reviewed or replaced with safer abstractions).

- Possibly some benchmarking to validate performance claims and optimize any slow paths.

- **Integration with Async Ecosystem (Long Term):** While BRRTRouter deliberately avoids Tokio/async for now, the Rust ecosystem is largely async-focused. In the long run, the project might offer adaptors to integrate with async code – for example, allowing an async function to be a handler (by spawning it in a May coroutine or bridging via `std::thread`). Another possibility is that Rust introduces first-class language support for coroutines (the `gen` keyword mentioned by the author ⁴⁴). If that happens, BRRTRouter could potentially migrate from May to the new native coroutine system, gaining broader support while retaining its model. This is speculative, but the mention shows the project's philosophy of leveraging Rust's evolving concurrency features.

- **Expanded Protocol Support:** Currently, BRRTRouter focuses on HTTP/1.1 (via `may_minhttp`). A future extension could be HTTP/2 or HTTP/3 support, or integration with WebSockets (which could be implemented by upgrading an HTTP connection and managing it in a coroutine). Another area

could be RPC support or bridging to gRPC by translating OpenAPI (or AsyncAPI) specs to handlers in a similar coroutine-driven manner.

- **Community and Contributors:** As the project stabilizes, it will benefit from more contributors and feedback. Writing more examples (for instance, an example with authentication, or file uploads, etc.) will help users adopt it. The **documentation site or book** will be important for driving adoption. The author has indicated a doc book is planned ⁴³. We may see the `docs/` directory in the repo get populated with usage guides, troubleshooting tips, and performance tuning advice.

Conclusion

BRRTRouter is an ambitious project that brings together **OpenAPI-driven development** and **Rust coroutine** concurrency. Its design favors a blocking coding style and dynamic reloading, which will feel familiar to developers coming from languages like Python (FastAPI) or JavaScript, while still harnessing Rust's speed and safety. The technical analysis of the test failures shows that most issues are understood and solvable with targeted fixes (panic handling, synchronization improvements). Once those are addressed, BRRTRouter should have a solid foundation for a 1.0 release.

In this report, we've provided a deep dive into how BRRTRouter works, identified the causes of its intermittent test failures, and suggested concrete improvements. By implementing these fixes and continuing to refine the architecture (especially around error handling and code generation), BRRTRouter can become a robust framework for building APIs in Rust without the complexity of `async/await`. The future roadmap is exciting – with full spec-driven codegen, hot-reload, and possibly new language features, BRRTRouter has the potential to significantly improve productivity for Rust web developers who want high performance *and* high ergonomics.

<c**Sources:** This analysis was based on the BRRTRouter codebase and commit history as of May 30, 2025. Key excerpts from the code and repository were cited to support the statements: repository description ¹, test code and comments ²⁹ ³³ ⁴⁵, and implementation details from commits (hot reload logic ²³, server startup ²², etc.). The Rust forum discussion by the project author provided additional context on goals and design philosophy ⁴³ ⁴⁴. These inline references point to the exact lines in the source or discussion for verification and further reading.]

¹ GitHub - microscaler/BRRTRouter: BRRTRouter is a high-performance, coroutine-powered request router for Rust, driven entirely by an OpenAPI 3.1.0 Specification

<https://github.com/microscaler/BRRTRouter>

² ³ ⁴ ⁵ ⁶ ⁷ ¹⁸ ²⁰ ²¹ ²² ²⁵ ²⁸ ³⁵ ³⁶ Fix SSE flag parsing and ensure tests set stack size (#94) · microscaler/BRRTRouter@4e33779 · GitHub

<https://github.com/microscaler/BRRTRouter/commit/4e3377923cadf5259d63dd9907c7044a5b6f2394>

⁸ ⁴⁰ ⁴¹ Add tracing to coroutine tests by casibbald · Pull Request #113 · microscaler/BRRTRouter · GitHub

<https://github.com/microscaler/BRRTRouter/pull/113>

⁹ ¹⁰ ¹¹ ¹² ¹³ ²⁴ Commits · microscaler/BRRTRouter · GitHub

<https://github.com/microscaler/BRRTRouter/commits/main/>

14 15 16 17 23 Add hot reload support (#81) · microscaler/BRRTRouter@c3ff7df · GitHub

<https://github.com/microscaler/BRRTRouter/commit/c3ff7df447e34fbdd953b60f600d98c5fba2a337>

19 34 Add tracing to coroutine tests (#113) · microscaler/BRRTRouter@ac96a2d · GitHub

<https://github.com/microscaler/BRRTRouter/commit/ac96a2d9d3f49a4830e3e6cab2780d583cf4ced8>

26 27 29 30 31 32 33 37 38 39 45 comments on tests and rebase · microscaler/BRRTRouter@fb302c1 ·

GitHub

<https://github.com/microscaler/BRRTRouter/commit/fb302c17d7477e31b34f299cbec1c52faf0a18aa>

42 43 44 Rust Alternative to FastAPI - The Rust Programming Language Forum

<https://users.rust-lang.org/t/rust-alternative-to-fastapi/129619>